

CSE512-Project-Phase2-Code-Template (v3.0, Oct 26)

Requirement

In Project Phase 2, you need to write two User Defined Functions ST_Contains and ST_Within in SparkSQL and use them to do four spatial queries:

- Range query: Use ST_Contains. Given a query rectangle R and a set of points P, find all the points within R.
- Range join query: Use ST_Contains. Given a set of Rectangles R and a set of Points S, find all (Point, Rectangle) pairs such that the point is within the rectangle.
- Distance query: Use ST_Within. Given a point location P and distance D in km, find all points that lie within a distance D from P
- Distance join query: Use ST_Within. Given a set of Points S1 and a set of Points S2 and a distance D in km, find all (s1, s2) pairs such that s1 is within a distance D from s2 (i.e., s1 belongs to S1 and s2 belongs to S2).

A Scala SparkSQL code template is given. You must start from the template.

The main function is in "SparkSQLExample.scala".

The User Defined Functions that you need to implement are in "SpatialQuery.scala". Replace the query part with your code from the old template.

The detailed requirements are as follows:

1. ST_Contains

Input: pointString:String, queryRectangle:String

Output: Boolean (true or false)

Definition: You first need to parse the pointString (e.g., "-88.331492,32.324142") and queryRectangle (e.g., "-155.940114,19.081331,-155.618917,19.5307") to a format that you are comfortable with. Then check whether the queryRectangle fully contains the point. Consider on-boundary point.

2. ST_Within

Input: pointString1:String, pointString2:String, distance:Double

Output: Boolean (true or false)

Definition: You first need to parse the pointString1 (e.g., "-88.331492,32.324142") and pointString2 (e.g., "-88.331492,32.324142") to a format that you are comfortable with. Then check whether the two points are within the given distance. Consider on-boundary point. To simplify the problem, please assume all coordinates are on a planar space and calculate their Euclidean distance.

3. Use Your UDF in SparkSQL

The code template has loaded the original data (point data, arealm.csv, and rectangle data, zcta510.csv) into DataFrame using tsv format. You don't need to worry about the loading phase.

Range query: `select * from point where ST_Contains(point._c0, '-155.940114,19.081331,-155.618917,19.5307')`

Range join query: `select * from rectangle,point where ST_Contains(rectangle._c0,point._c0)`

Distance query: `select * from point where ST_Within(point._c0, '-88.331492,32.324142',10)`

Distance join query: `select * from point p1, point p2 where ST_Within(p1._c0, p2._c0, 10)`

4. Run your code on Apache Spark using "spark-submit"

If you are using the Scala template, note that:

1. You **only have to replace the logic** (currently is "true") in all User Defined Function.
2. The main function in this template takes **dynamic length of parameters** as follows:
 - Output file path (**Mandatory**): `/Users/ubuntu/Downloads/output`
 - Range query data file path, query window: `rangequery /Users/ubuntu/Downloads/arealm.csv -155.940114,19.081331,-155.618917,19.5307`
 - Range join query data file path, range join query window data file path:
`rangejoinquery /Users/ubuntu/Downloads/arealm.csv /Users/ubuntu/Downloads/zcta510.csv`

- Distance query data file path, query point, distance: `distancequery /Users/ubuntu/Downloads/arealm.csv -88.331492,32.324142 10`
- Distance join query data A file path, distance join query data B file path, distance: `distancejoinquery /Users/ubuntu/Downloads/arealm.csv /Users/ubuntu/Downloads/arealm.csv 10`

- The number of queries and the order of queries in the input **do not matter**. The code template will detect the corresponding query and call it!
- Two example datasets are put in "src/resources" folder. arealm10000 is a point dataset and zcta10000 is a rectangle dataset. You can use them to test your code but eventually you must run your code on NYC taxi trip dataset. Our auto-grading system will also run your code on many different datasets.
- Here is an example that tells you how to submit your jar using "spark-submit" `./bin/spark-submit ~/GitHub/CSE512-Project-Phase2-Template/target/scala-2.11/CSE512-Project-Phase2-Template-assembly-0.1.0.jar ~/Downloads/`
- A test case file is given: `exampleinput` . A correct answer is given: `exampleanswer`

5. Vocareum (submission site)

In Vocareum, all groups share a single but powerful cluster.

Vocareum is for submission only. After uploading your jar, directly click "submit".

Do not test your code here. You have limited submission times. Every submission impacts your grade!

Website

<https://www.vocareum.com/>

Your Spark job log is in "aggregated.log" file of your Vocareum "LatestSubmission" folder

6. Thoth Lab (testing environment, optional)

In Thoth Lab, two groups share an individual cluster that has 3 VMs. Feel free to use this cluster as your testing environment. Any operations on Thoth Lab cluster do not impact your grade.

File structure

On all machines in your cluster:

`/CSE512/data/nyc-spark/pickup/` has all point data stored in `x,y` CSV format.

`/CSE512/code/hadoop-2.7.4` has a fresh Hadoop 2.7.4 `/CSE512/code/spark-2.2.0-bin-hadoop2.7` has a fresh Spark 2.2.

You need to set up password-less SSH and start HDFS/Spark by yourself. These machines have static IP.

Website

User credential activate website (for activation only!): <https://www.thothlab.org/>

Thothlab Cluster Interface: <https://openstack.thothlab.com>

Domain:users

Submission

Deadline

November 2nd, 2017 11:59 pm

Submission files

- Submit your code on Vocareum website in order to get your grade.
- Submit your project source code onto Blackboard in a compress zip file of "cse512-phase2-GROUPNAME" for plagiarism detection. Note that: you need to make sure your code can compile and package by entering `sbt clean assembly` . We will run the compiled package on our cluster directly using "spark-submit".
- If your code cannot compile and package, you will not receive any points.

How to debug your code in IDE

If you are using the Scala template

1. Use IntelliJ Idea with Scala plug-in or any other Scala IDE.
2. Replace the logic of User Defined Functions ST_Contains and ST_Within in SpatialQuery.scala.
3. Append `.master("local[*]")` after `.config("spark.some.config.option", "some-value")` to tell IDE the master IP is localhost.
4. In some cases, you may need to go to "build.sbt" file and change `% "provided"` to `% "compile"` in order to debug your code in IDE
5. Run your code in IDE
6. **You must revert Step 3 and 4 above and recompile your code before use spark-submit!!!**

How to submit your code to Spark

If you are using the Scala template

1. Go to project root folder
2. Run `sbt clean assembly` . You may need to install sbt in order to run this command.
3. Find the packaged jar in `"/target/scala-2.11/CSE512-Project-Phase2-Template-assembly-0.1.0.jar"`
4. Submit the jar to Spark using Spark command `"/bin/spark-submit"`